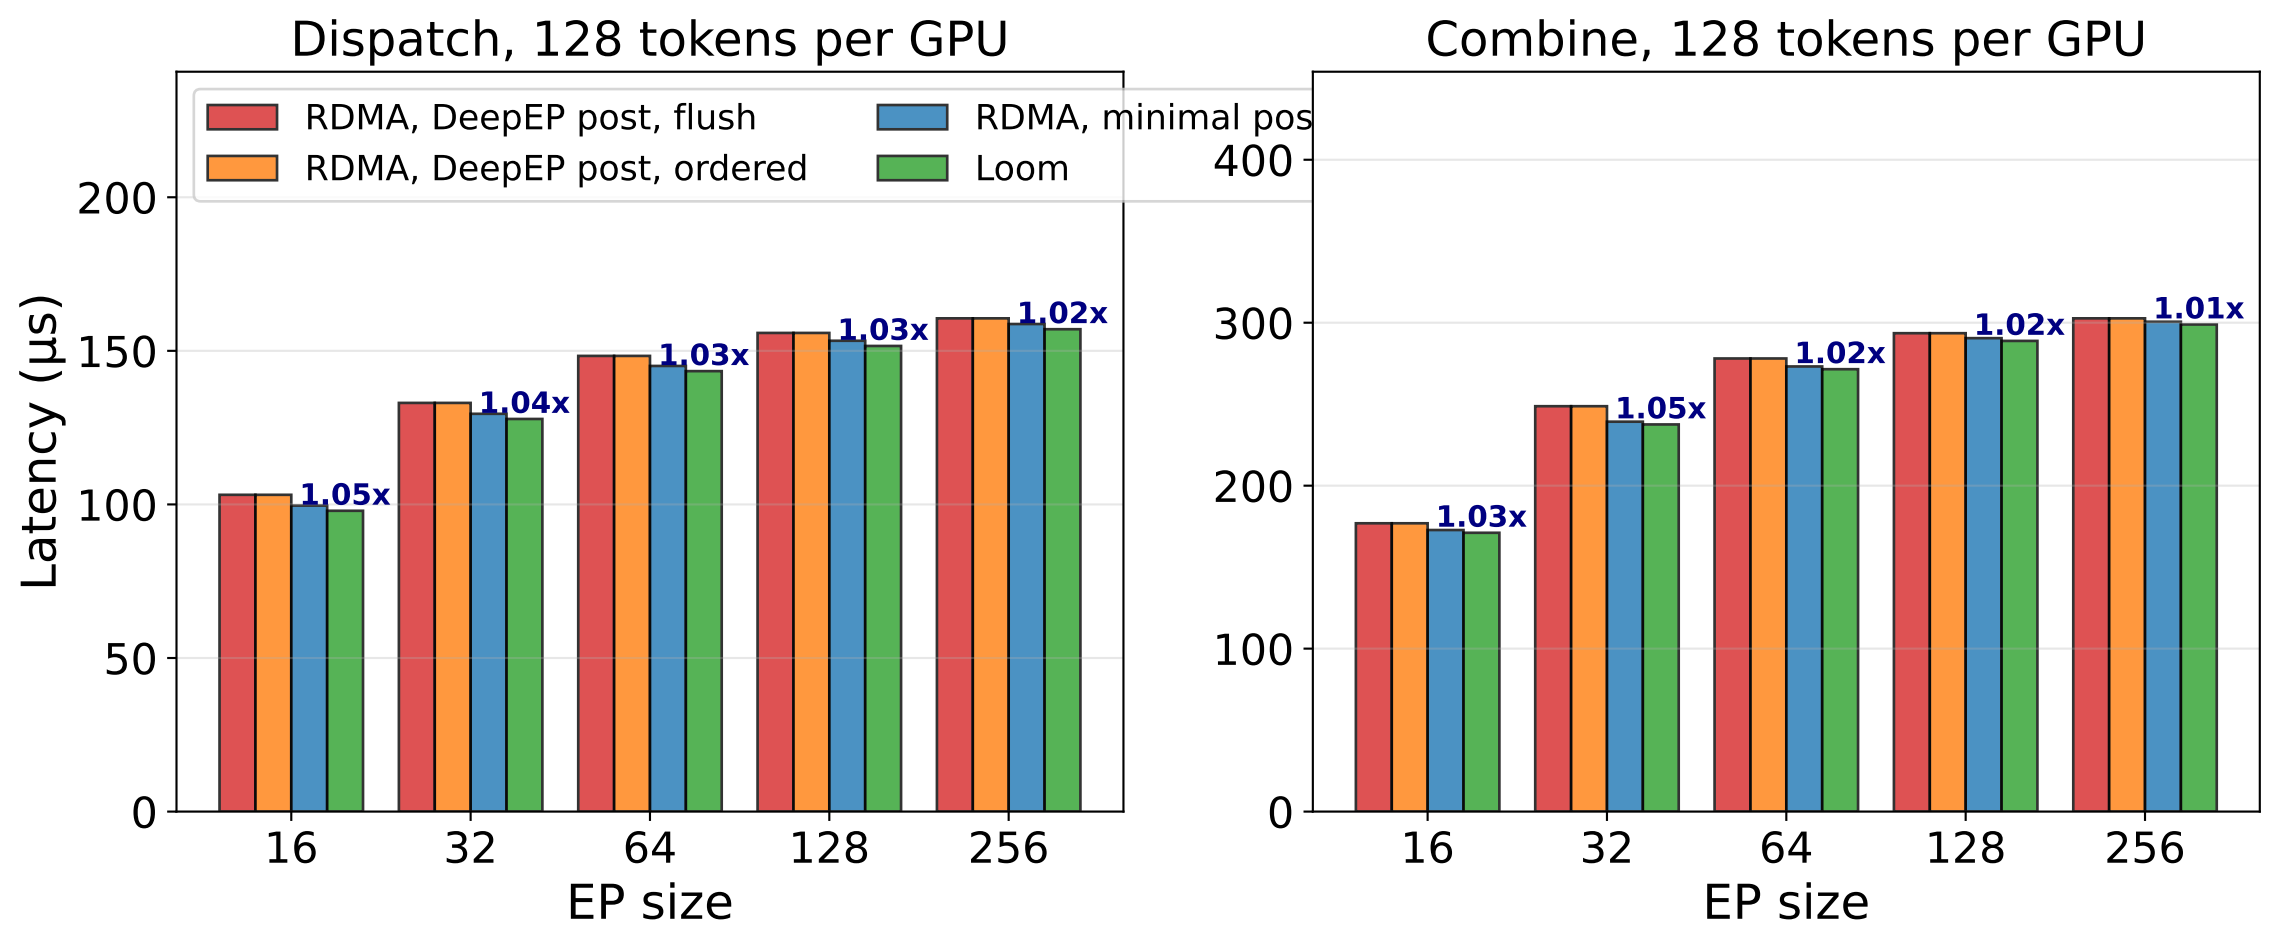
Dispatch, 128 tokens per GPU
RDMA, DeepEP post, flush
RDMA, DeepEP post, ordered
RDMA, minimal pos
Loom
Latency (µs)
0
50
100
150
200
1.05x
1.04x
1.03x
1.03x
1.02x
16
32
64
128
256
EP size
Combine, 128 tokens per GPU
0
100
200
300
400
1.03x
1.05x
1.02x
1.02x
1.01x
16
32
64
128
256
EP size