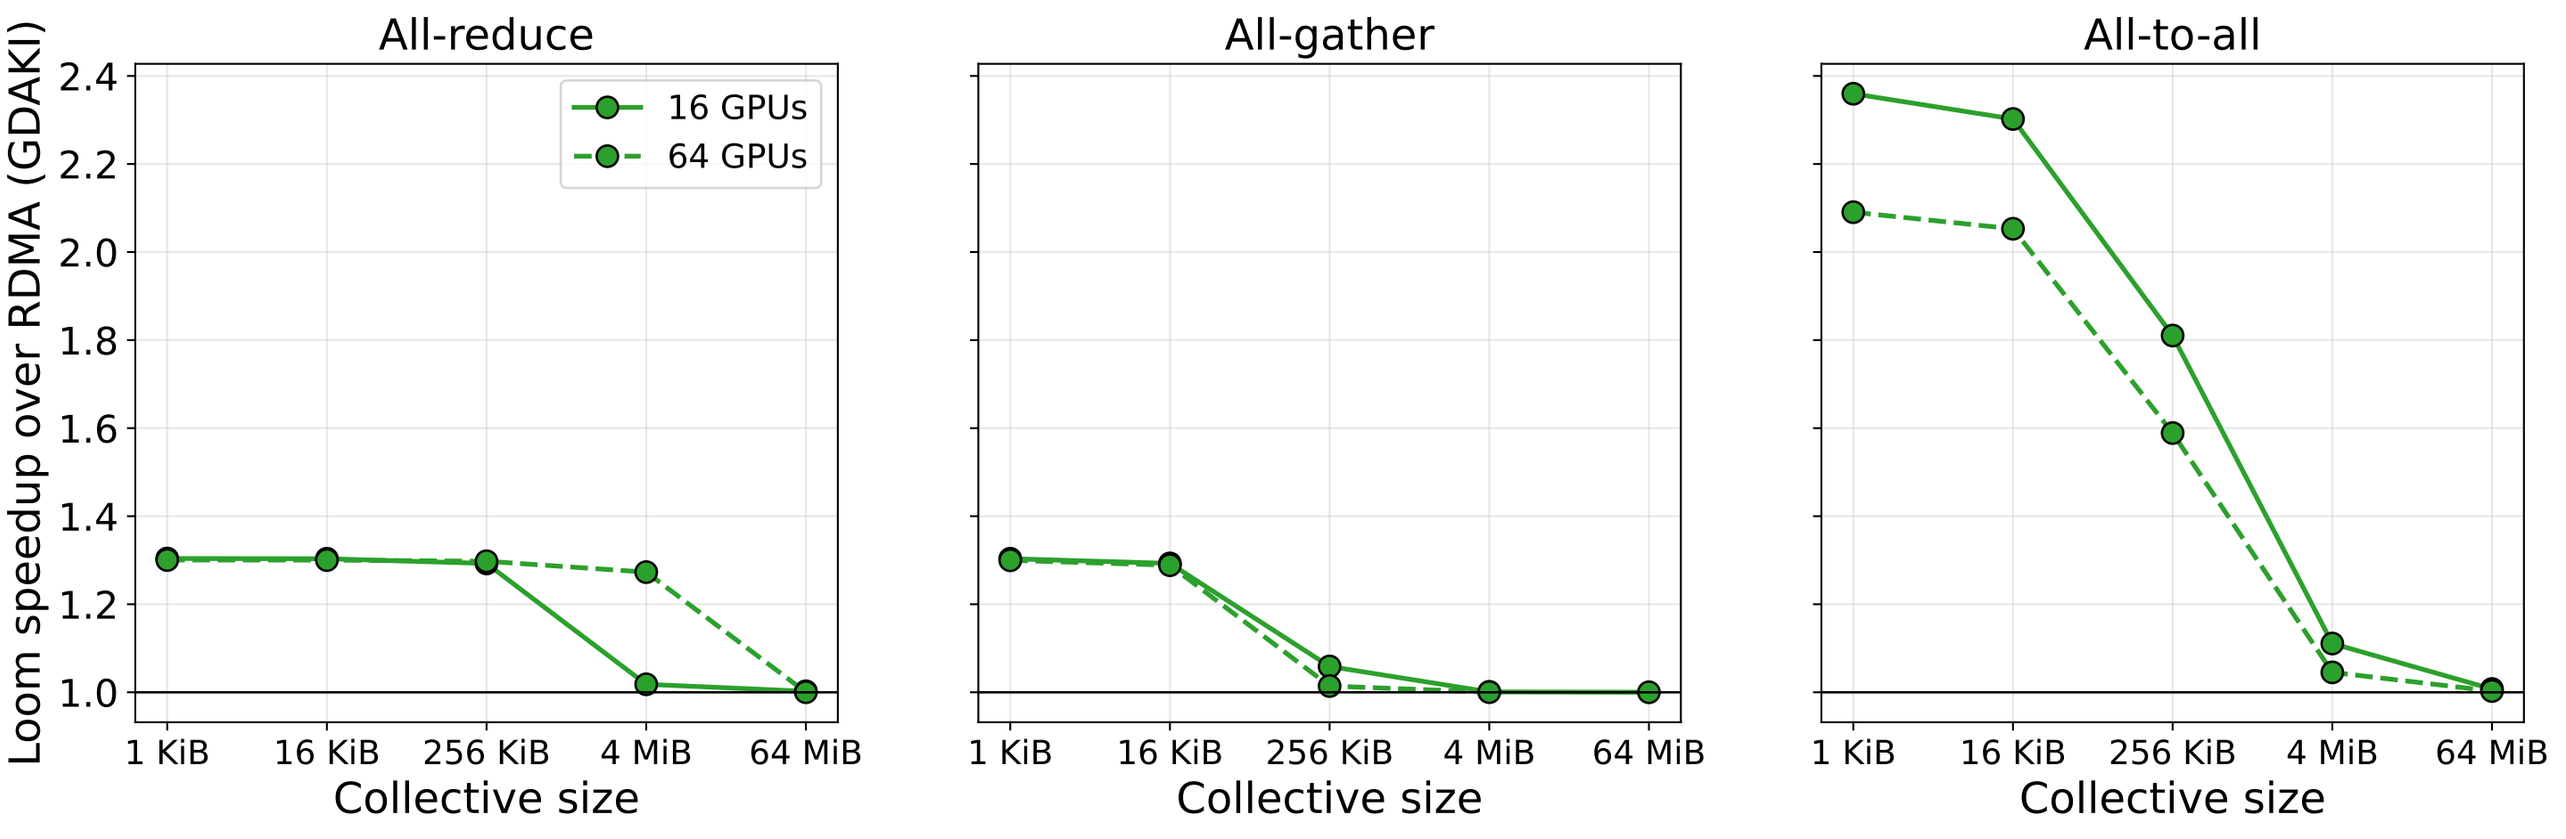
All-reduce
All-gather
All-to-all
Loom speedup over RDMA (GDAKI)
2.4
2.2
2.0
1.8
1.6
1.4
1.2
1.0
16 GPUs
64 GPUs
1 KiB
16 KiB
256 KiB
4 MiB
64 MiB
Collective size