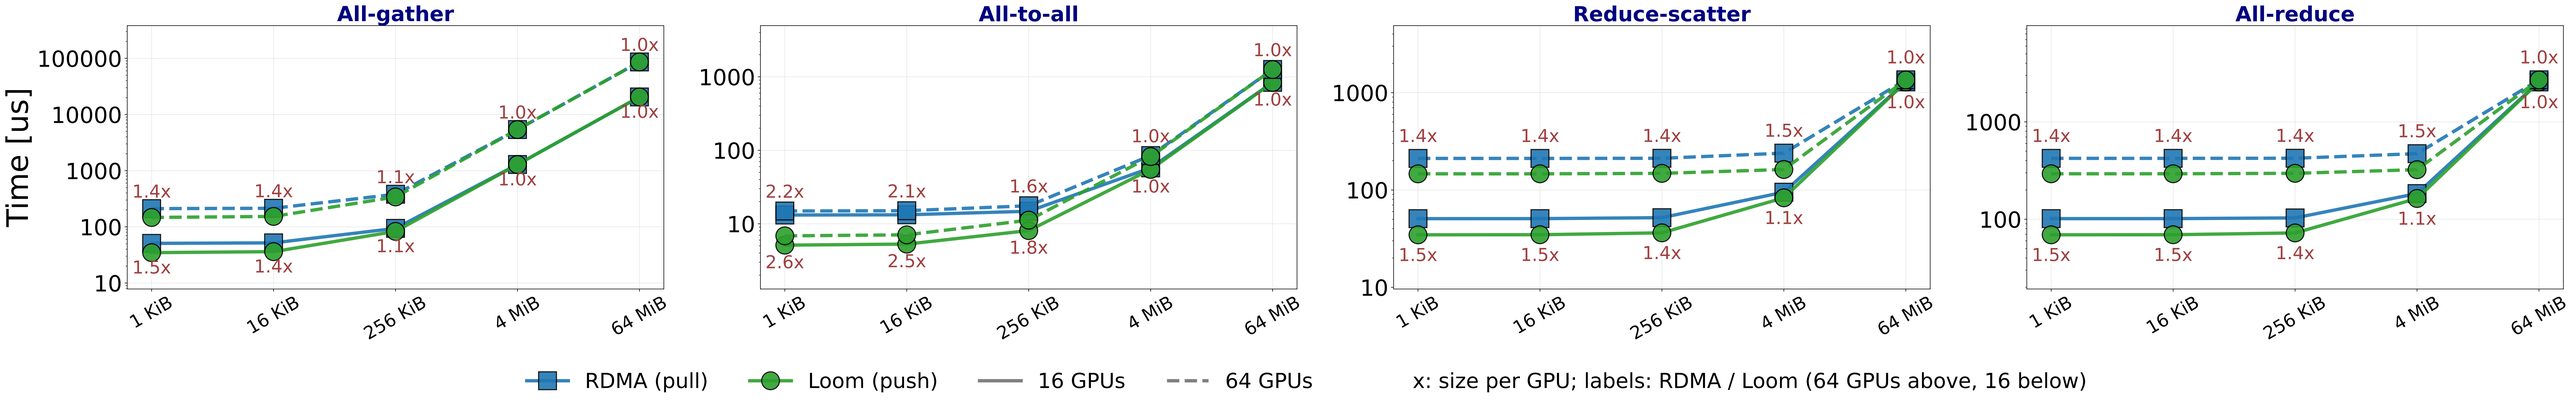
All-gather
All-to-all
Reduce-scatter
All-reduce
Time [us]
100000
10000
1000
100
10
1 KiB
16 KiB
256 KiB
4 MiB
64 MiB
1.4x
1.5x
1.1x
1.0x
2.2x
2.6x
2.1x
2.5x
1.6x
1.8x
RDMA (pull)
Loom (push)
16 GPUs
64 GPUs
x: size per GPU; labels: RDMA / Loom (64 GPUs above, 16 below)