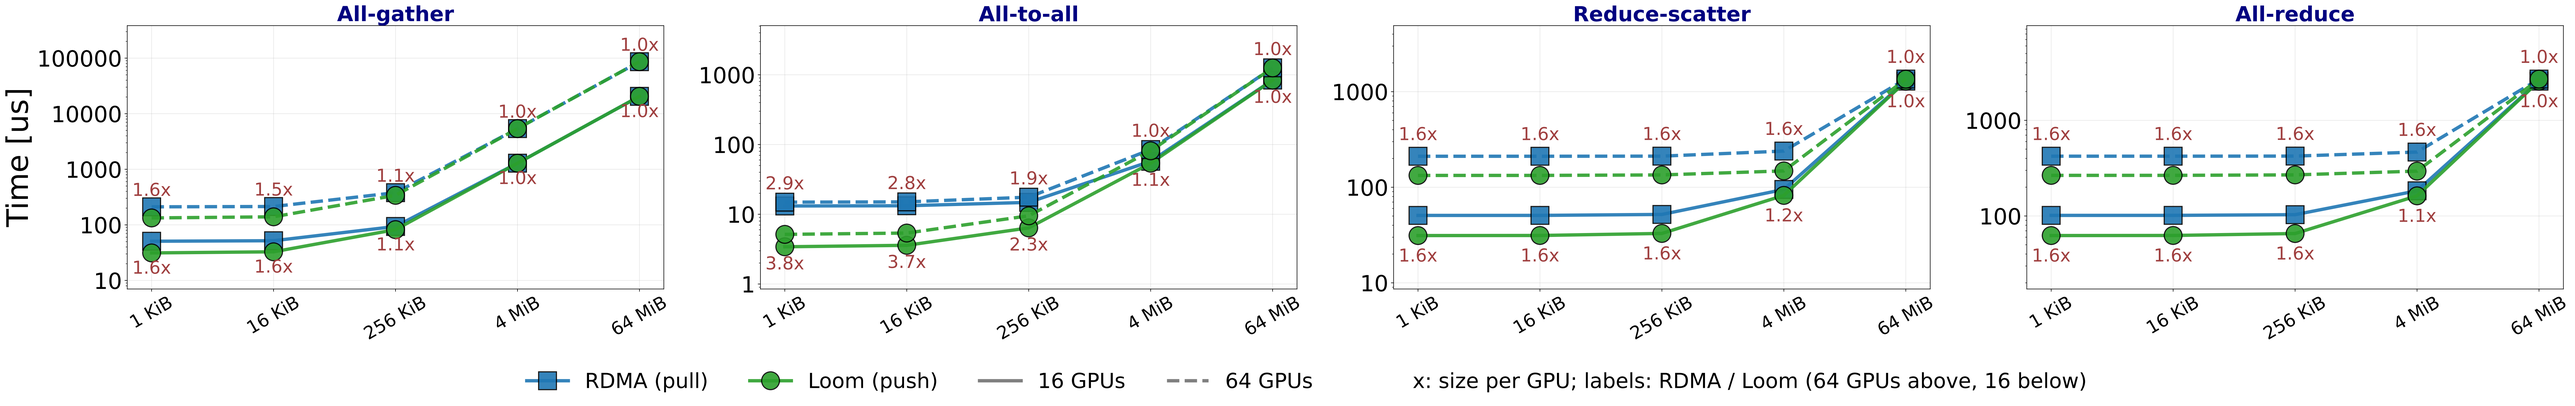
All-gather
All-to-all
Reduce-scatter
All-reduce
Time [us]
100000
10000
1000
100
10
1
1 KiB
16 KiB
256 KiB
4 MiB
64 MiB
1.6x
1.5x
1.1x
1.0x
2.9x
2.8x
1.9x
3.8x
3.7x
2.3x
RDMA (pull)
Loom (push)
16 GPUs
64 GPUs
x: size per GPU; labels: RDMA / Loom (64 GPUs above, 16 below)