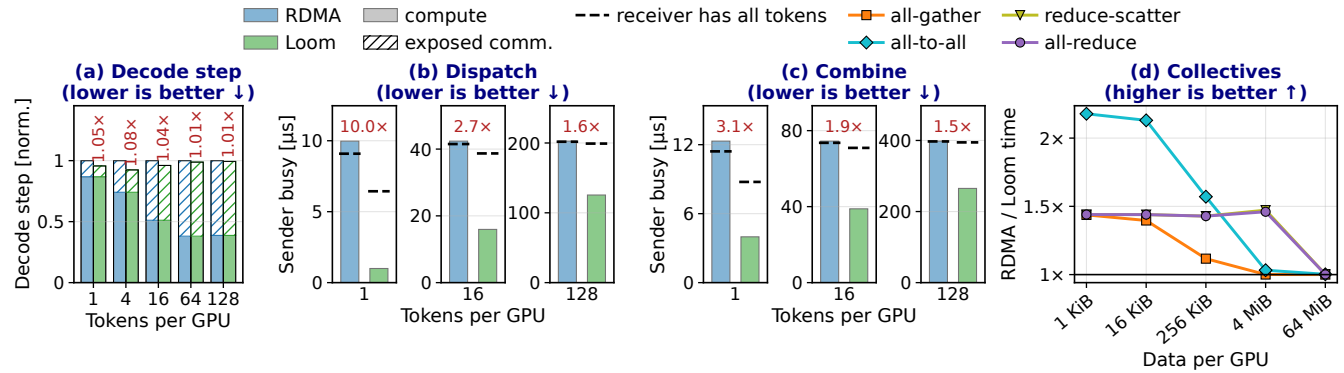
RDMA
Loom
compute
exposed comm.
receiver has all tokens
all-gather
all-to-all
reduce-scatter
all-reduce
(a) Decode step
(lower is better ↓)
Decode step [norm.]
1.05×
1.08×
1.04×
1.01×
1.01×
Tokens per GPU
(b) Dispatch
(lower is better ↓)
Sender busy [µs]
10.0×
2.7×
1.6×
Tokens per GPU
(c) Combine
(lower is better ↓)
Sender busy [µs]
3.1×
1.9×
1.5×
Tokens per GPU
(d) Collectives
(higher is better ↑)
RDMA / Loom time
Data per GPU