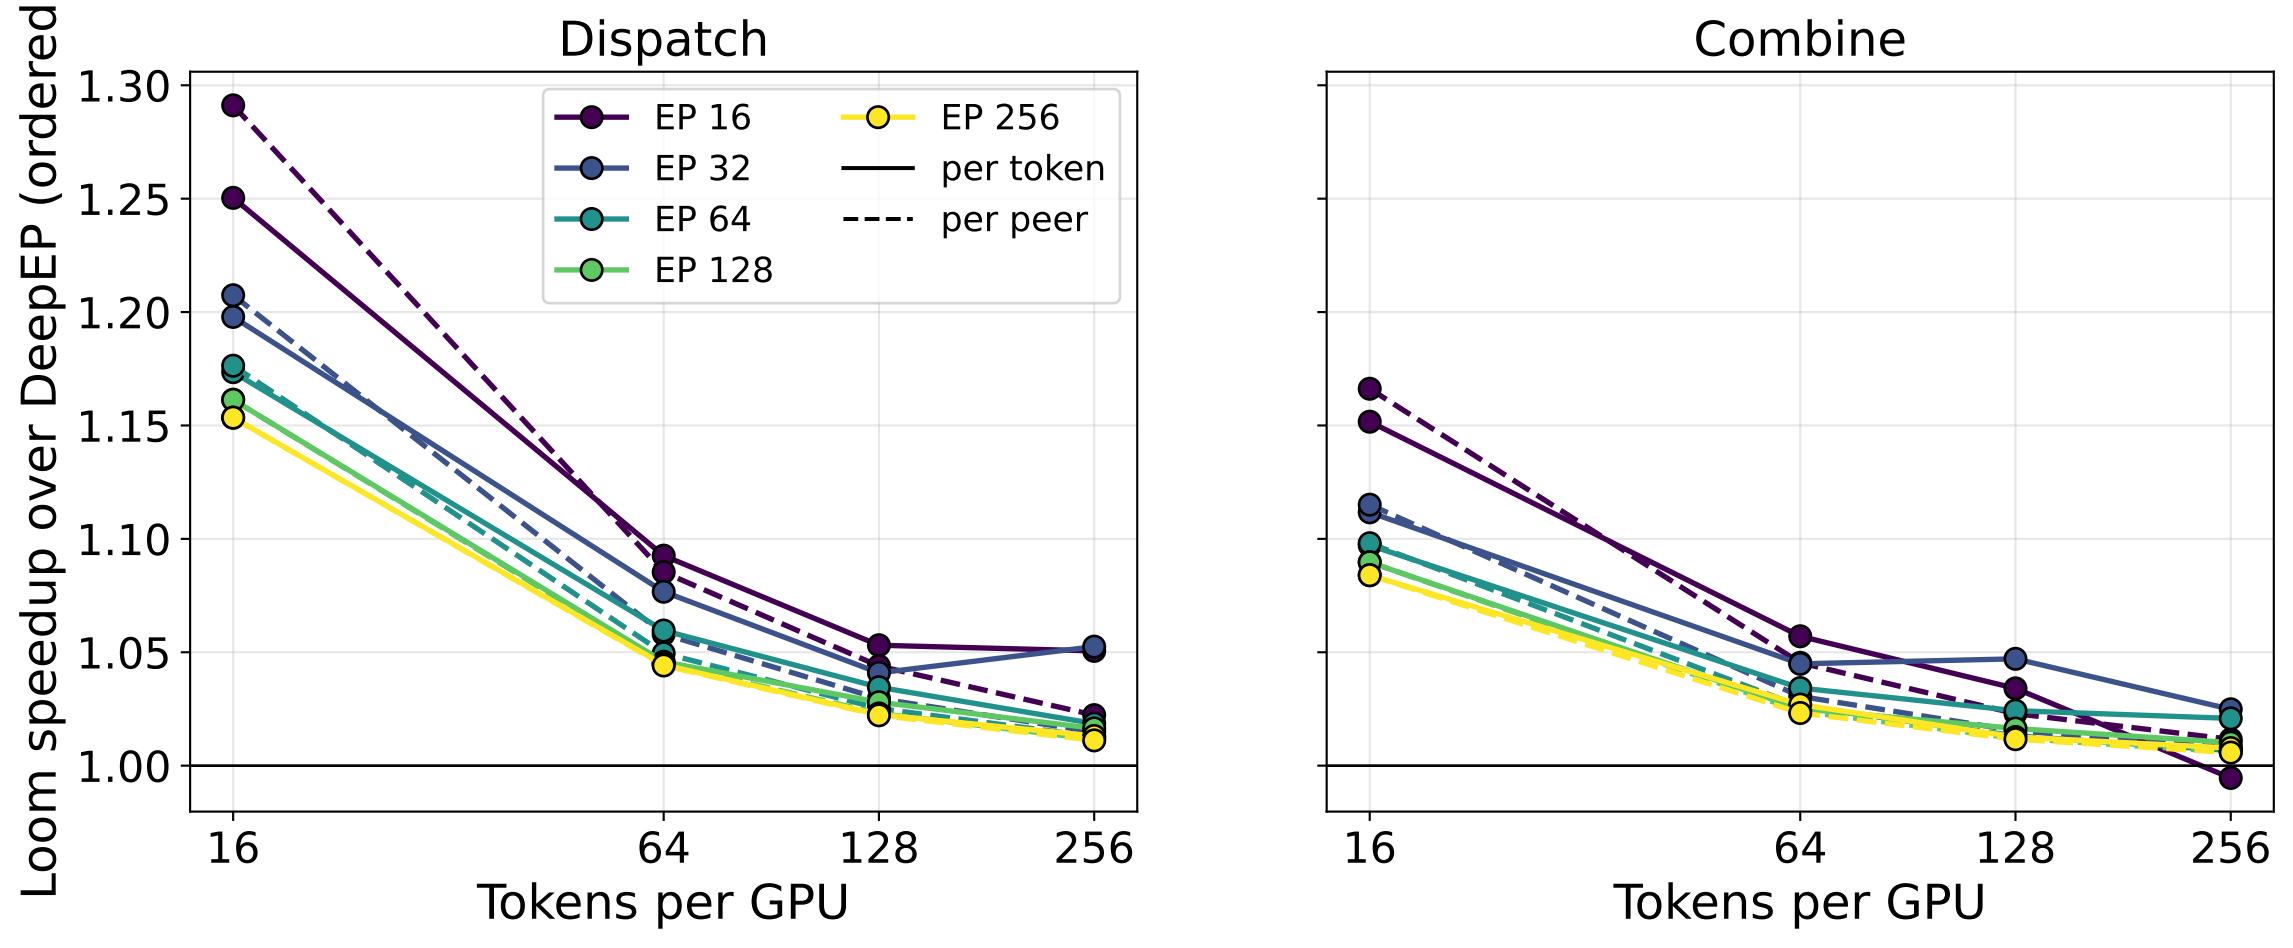
Dispatch
Combine
Loom speedup over DeepEP (ordered
1.30
1.25
1.20
1.15
1.10
1.05
1.00
16
64
128
256
Tokens per GPU
EP 16
EP 32
EP 64
EP 128
EP 256
per token
per peer